

# Non-Parametric Analysis of Social Influence Bias in the California Report Card

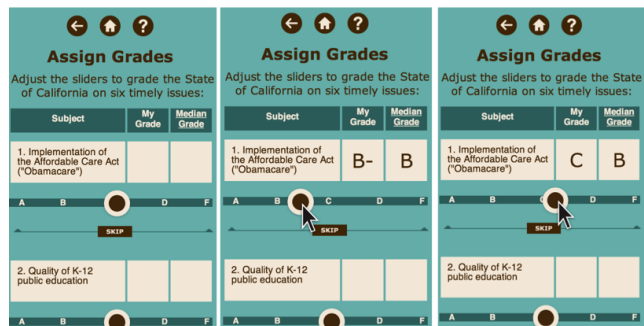
## ABSTRACT TODO

## 1. INTRODUCTION

User-assigned ratings are a key component of almost all recommender systems. Rating data, like traditional surveys, are subject to a variety of biasing tendencies [11]. We explore a class of biases, collectively called *Social Influence Bias*, which arise from feedback from the actions of other users in the system [9, 15, 21]. One aspect of Social Influence Bias is the phenomenon called *Social Herding*, where the feedback from the community encourages future participants to conform to what they perceive as the “norm” in the community. In a rating system, this can lead to an increased tendency to leave ratings close to the mean or the median rating. The effects of social herding are crucial to the design of recommendation algorithms as many algorithms assume statistical independence between different users and use the spatial relationships between numerical features representing those users.

A common feedback mechanism is the use of aggregate statistics, for example, showing the average rating for a product before a participant shares his or her rating (Figure 1). In many cases, such as product reviews, it is not practical to hide this information from potential raters. The use of social content is an established user experience design technique to incentivize participation and increase user engagement with the tool [19]. Furthermore, an application of particular interest is online participatory democracy where open aggregate results increase the transparency of the system [1, 18, 17].

In recent related work, Muchnik et al. [16], used a randomized experiment to determine the magnitude of social herding in up-voting in Reddit.com. They randomly treated forum posts with extra up-votes and down-votes and measured the treatment effect; concluding that a statistically significant social herding tendency exists. We study a re-



**Figure 1: Grading in the California Report Card.** Participants enter grades on six timely issues facing the State of California. After entering their grades, the median grade over all participants is revealed. Participants have the option to change their grades after seeing the median. We model the tendency to regress towards the medians.

lated effect using a new platform, the California Report Card (CRC), which reveals the median values to users *after* they assign their own ratings, and then allows them to modify their ratings. We test whether users who have already submitted a rating will actually *change* their existing ratings upon learning the median rating for the population and if there is a tendency to “herd” towards the median grade.

For comparison with the CRC, we ran a reference survey through SurveyMonkey which was given to a random sample (N=611) from the company’s paid pool of California participants. In this survey, we asked users to respond to the same questions as the CRC on the same grading scale; but without the feedback of the median grade.

The findings of Muchnik et al. suggest that social herding will be observed in the form of a regression towards the observed median grade during the grade changes. In this paper, we test the hypothesis of social herding, propose a model for the relationship between observed medians and subsequent grade change, and provide experimental results comparing the data from the CRC to the randomized reference survey.

## 1.1 Hypotheses and Contributions

**Null Hypothesis** Viewing the median grade does not affect how a user chooses to change his or her grade, and does not

affect any future grades given by the participant.

**Social Herding Towards the Median** Grade changes are on average towards the median grade. Also, the final grades of participants who change their grades are more tightly concentrated around the median from participants who did not change their grades and participants from the reference survey.

**Social Herding and Question Order** Disagreement with the median on previous questions affect how a participant grades future questions. Participants who disagreed with the median greatly are more likely to leave future responses that are closer to the median.

We develop non-parametric testing procedure, based on the Wilcoxon Rank-Sum statistic (also called the Mann-Whitney statistic) [13], to test these hypotheses. We chose a non-parametric framework because Muchnik et al. focused on only a binary input mechanism (up or down vote) and we extended this analysis to grading sliders with 13 possible values from (A+ to F) without having to make strong assumptions about the distribution of grades. In addition to the hypothesis testing, we model grade changes with a polynomial regression. As before, to avoid having to make a strong assumption about the structure of the model, we use an information theoretic model to learn a flexible degree polynomial.

## 2. RELATED WORK

In Asch’s famous conformity experiments [3, 2, 6], groups of participants were asked to match a line with a set of three different sized lines one of which was of the correct size. In reality, only one of the participants was “real” and the others were actors who unanimously chose an incorrect choice. On average, 25% of participants conformed to the incorrect consensus compared to 1% of incorrect answers in a control group.

The Asch model for conformity is the theoretical basis for social herding [4, 5], and herding has been a popular consumer choice model in economics [7, 10, 12]. Such models have also been studied in psychology as “persuasion bias” [9]. In 2011, Lorenz et al. described how these biases can undermine the effectiveness of crowd intelligence in estimation tasks [14]. They argue that social herding causes a diminished diversity of opinion potentially leading to inefficiencies and inaccurate collective estimates. Danescu-Niculescu-Mizil et al. analyze helpfulness ratings on Amazon product reviews [8]. They found that the helpfulness ratings did not just depend on the content of the review but also its aggregate score and its relationship to other scores. In order to better distinguish social influence from other biases, Muchnik et al. designed a randomized experiment in which comments on Reddit.com were randomly up-treated or down-treated [16]. They concluded a statistically significant bias where a positive treatment increased the likelihood of positive ratings by 32%. In both Danescu-Niculescu-Mizil et al. and Muchnik et al., they looked at the problem of Social Influence bias in an a priori setting, where users see the aggregate statistic before giving their rating. Our work tests for a particular form of social influence where users are given the opportunity to change their opinions following the feedback.

Zhu et al. conducted an experiment in which users evaluate an image on a subjective question with binary scale (eg. “Is this image cute?”), which was followed (either immediately or later) by a presentation of the crowd consensus opinion [22]. Users were given an opportunity to change their response, and they concluded that there was a significant tendency to change submissions. The tendency to change was the strongest when users were asked to make their second decision much later and not immediately after the first. However, Zhu et al. also acknowledge there are competing psychological factors at work in this experiment. Along these lines, Sipos et al. argue that context along with an aggregate rating plays a large role in the users’ ratings. That is, users may attempt to “correct” the average, by voting in a more polarizing manner (more positively or negatively) [20]. We extend this prior work to measure and predict these changes when the input is more complex than a binary scale, and propose a non-parametric methodology that can be, in principle, extended to a variety of different input mechanisms. Our model can also account for a changing aggregate statistic such as a median rating changing as more data is collected.

## 3. THE CALIFORNIA REPORT CARD

### 3.1 System Description

The California Report Card (CRC) is a web application that allows participants to advise the state government on timely policy issues. When participants arrive at the application, they “grade” the state on following six issues:

- Implementation of the Affordable Care Act (“Obamacare”)
- Quality of K-12 public education
- Affordability of state colleges and universities
- Access to state services for undocumented immigrants
- Laws and regulations regarding recreational marijuana
- Marriage rights for same-sex partners

Grades were assigned on a thirteen point scale (A+,A,A-,...,D-,F). These issues were posed in a fixed sequential order each with the same input scale. Participants submitted grades using a click-and-drag slider interface as illustrated in Figure 1. On mobile devices this slider required the participants to touch and drag their finger to the desired grade.

Upon release of the slider, the CRC reveals the median grade for that issue over all prior participants. Even after the median grade is revealed the slider is still active and participants can change their grades. However, it is important to note that participants were not explicitly told that they could change their grades. Another important observation is that participants who accessed the application at different times may have seen different median grades as they were calculated based on the data up to that point. We recorded the initial grade, the median that the participant observed, and any subsequent changes along with timestamps for each of the events. Grading all of the six issues was not mandatory and participants had the option to skip any of the issues.

The CRC has an additional open-ended discussion phase where participants submit textual suggestions on future issues to include in the report card. In this work, we focus on the first phase and defer an analysis of biases in the discussion phase to future work.

### 3.2 Notation

To analyze this data, we mapped these 13 grades onto a scale from 0 to 1, with 1 being an A+ and 0 being an F. Let  $P$  denote the set of all participants. For each participant  $p_j \in P$ , we associate a 3-tuple of grades  $(g_i[j], m[j], g_f[j])$  which represent the initial grade, median observed by the participant, and the final grade. For each issue, we divided the participants into three subsets of  $P$ : ones who did not change their grades  $P_n$ , ones who changed  $P_c$ , and ones who skipped the question  $P_s$ . Our primary objective is to test the distributional properties of rating tuples from participants in  $P_n$  compared to those in  $P_c$ .

To ensure that all participants in the set  $P_c$  had an opportunity to see the median grade and then react, we filtered this group using the timestamps. The median grade appears in the interface with an animation whose completion time varied between devices, so we set a grace period of 3 seconds before we categorized the participant into set  $P_c$ .

For consistency, we use the same notation to describe participants in the reference survey. We denote the set of reference survey participants as set  $R$ , and each participant is associated with a 3-tuple  $(g_i[j], m[j], g_f[j])$ . However, since the reference survey does not reveal the aggregate statistics  $g_i[j] = g_f[j]$  and  $m[j]$  is the median of the prior participants (which is not shown).




Figure 2: Breakdown of Activity in CRC Assessment.

## 4. SOCIAL HERDING TOWARDS THE MEDIAN HYPOTHESIS TEST

There are three principle challenges in testing **Social Herding Towards the Median** hypothesis. The first challenge is that parametric significance tests comparing two sample means such as the two sample t-test and z-test are known to perform poorly for multimodal distributions. Another significance test that is commonly applied to compare spreads of distributions is the F-test, which is also known to perform poorly for many non-normal distributions [?]. Furthermore, this test is usually used to test the spread of data around

the mean, which only in very special conditions such as normal distributions, aligns with the median. The discreteness of our data leads to multi-modal distributions which are not optimal for these testing methods.

The second challenge is that there is a natural tendency for grades to concentrate around the median even without a bias. Consider the following participant behavioral model. Suppose that participants are not accustomed to a slider-based input. We can model the first grade that the participant leaves as uniformly randomly anywhere on the slider. As the participant begins to understand how to use the slider their use becomes more accurate, ultimately settling on a grade from our observed distribution of final grades. This model, the first grade is uniformly random and the second grade is a sample from the observed distribution, would result in a strong regression towards the median; even if there is no causal link with seeing the median.

Finally, the median grade  $m_i$  can be different for each participant. The median grade is calculated over all prior participants and thus is dependent on when the participant submitted their first grade. In practice, the median will eventually converge for a large number of participants, but it would be incorrect to measure concentration around a final median.

To address these three challenges, we propose a non-parametric model based on the Wilcoxon statistic to test the hypothesis that the group of participants that changed their grades are more tightly centered around the median grade that the participants observed. Our tests compare absolute deviations around the median for  $P_n$ ,  $P_c$ , and  $R$ ; which, as a relative test, controls for the natural tendency for grades group around the median. Furthermore, it is more robust to the effects of alternate models such as the one described in our second challenge in comparison to a direct test of correlation.

### 4.1 Non-parametric Significance Test

Recall that  $P_n$  is the set of participants that did not change their grades and  $P_c$  be the set of participants that changed their grades. We define a set  $X_c, X_n$  of absolute deviations from the observed median of the final grade for each group:

$$X_c = \{|m[j] - g_f[j]|\} \forall j \in P_c \quad (1)$$

$$X_n = \{|m[j] - g_f[j]|\} \forall j \in P_n \quad (2)$$

For the purposes of hypothesis testing herding behavior, we ignore the sign of the deviation. However, in Section 6, where we build a predictive model for the changes, we include the sign.

Now, for the set  $X_c$ , we calculate the Wilcoxon rank-sum statistic. We assign a rank to each of the absolute deviations in the union set  $\mathbf{X} = X_c \cup X_n$  (ie. the largest change has rank 1 and the smallest has rank  $|X_c \cup X_n|$ ). For  $X_c$ , we sum the ranks of the deviations within its set:

$$W_c = \sum_{j \in P_c} R_j \quad (3)$$

The **Null Hypothesis** is that absolute deviations in  $X_c$  are the same size as  $X_n$ . Under this null hypothesis  $median(X_n) =$

$median(X_c)$ , the ranks will be evenly distributed between each group. Therefore, the null expected value and variance of  $W$  is:

$$\mathbb{E}(W) = \frac{(|\mathbf{X}| + 1) \cdot |X_c|}{2} \quad (4)$$

$$var(W) = \frac{(|\mathbf{X}| + 1) \cdot |X_c| \cdot |X_n|}{12} \quad (5)$$

For the significance level  $\alpha$ , we can test the probability that our calculated  $W_c$  comes from the null distribution. In other words, the test calculates the probability that a random subset of users (ignoring the categorization  $P_n$  and  $P_c$ ) can have the observed difference in rank-sum values. A significant result means that for the participants that changed their grades the changed changes are more tightly centered around the median grade they observed.

The same analysis can be used to test  $X_c$  against the absolute deviations in the reference survey  $X_r$

$$X_r = \{|m[j] - g_i[j]|\} \forall j \in R \quad (6)$$

or for initial vs. final ratings in the change group  $X'_c$ :

$$X'_c = \{|m[j] - g_i[j]|\} \forall j \in P_c \quad (7)$$

## 5. SEQUENCE DEPENDENCE

In the CRC, we pose each of the six issues in a constant order. In this section, we develop a model for testing the effect of the sequence of ratings. In [?], the authors suggest that *path dependence* is significant in online tools. In particular, we test to see whether deviation from the median on previous issues is correlated with deviation on the current issue. We test the following hypothesis: if participants observe that their grades significantly deviate from the median on previous questions, their future responses will be more tightly centered around the median.

This hypothesis is challenging to test as responses to issues may be correlated; even excluding the bias. Consider the following example, if the grades are positively correlated, then low grades on one question could imply even lower grades on another. In this case, we would see an increase in deviations even though it is not attributable to the biasing tendency. Consequently, we build a model that compares the CRC to the SurveyMonkey reference survey. We test to see if the relationship between the deviation of a participant's past grades and their current grades is different between the CRC and reference survey.

Let  $d_{ij}$  be the absolute deviation from the median grade of participant  $j$ 's grade on issue  $i$ . We define a statistic  $P_{ij}$ , which is the mean of all of the absolute deviations on the previous issues:

$$P_{ij} = \frac{1}{i-1} \sum_{k < i} d_{kj} \quad (8)$$

For each issue  $i > 1$ , we can get a set of differences between the absolute deviation of the current issue and the average previous absolute deviations:

$$D = \{(P_{ij} - d_{ij})\} \forall j \quad (9)$$

We can calculate the same statistic  $D_r$  for deviations from the reference survey. For a given issue, these two sets illustrate

the trend in deviations from the median. A large positive value implies that a participant who disagreed greatly with the median grade before is now much closer to the median. Conversely, a negative value implies their response deviates more.

While this statistic is difficult to interpret for an individual participant as their assessments may vary issue to issue, we can compare the distributions of differences from the CRC and Reference Survey. The two sets can be tested with the Wilcoxon model in the previous section. The results in [?] suggest that the CRC should show larger differences; corresponding to increasingly moderate grades by participants who observed that they disagreed with the median in the past. Thus, we test to see if the differences in the set from the CRC  $D$  are statistically significantly higher than those from the reference survey  $D_r$ . The Wilcoxon testing procedure is the following: (1) we rank the differences in  $D \cup D_r$ , (2) we calculate  $W$  which is the sum of the ranks in  $D$ , and (3) using the equation from the previous section we test the calculated  $W$  under the null hypothesis distribution.

A significant result means that in comparison to the reference survey, CRC participants future responses were more concentrated around the median (ie. a higher difference between  $P_{ij} - d_{ij}$ ). This test is particularly interesting in the context of initial grades rather than final ones. We can test to see how the concentration of grades around the median changes even without the biasing effect of revealing the median, and whether participants have a tendency to *guess* the median grade.

## 6. PARAMETER ESTIMATION FOR GRADE CHANGE MODEL

In the previous sections, we proposed a technique to test the significance of the regression towards the median. In this section, we build a model to describe the relationship between the variables in the 3-tuple  $(g_i[j], m[j], g_f[j])$ . We also contrasted two different parameters of interest: correlation and absolute deviation. In this section, we will further build on this to estimate two quantities: a functional relationship between  $m[j] - g_i[j]$  and  $g_f[j] - g_i[j]$ , and a quantification of how much more concentrated the changed grades are  $\Delta$ . The functional relationship, related to the correlation, will tell us how to predict a final grade given an observed median. The  $\Delta$  parameter will tell us how much more tightly grouped around the median the final grades are.

### 6.1 Modeling Heterogenous Changes

Previous work, suggests that Social Influence bias is not homogenous; that is a negative influence is different in magnitude than a positive influence [?]. This means that we cannot assume that the relationship between  $m[j] - g_i[j]$  and  $g_f[j] - g_i[j]$  is linear.

Similar to the previous section where we applied non-parametric tests, we propose a information theoretic model search that allows flexible parameter selection without making strong assumptions about the nature of the relationship. Let  $f \in \mathcal{P}^k$  be a polynomial of degree  $k$ . The square loss of  $f$ , is the

error in predicting  $g_f[j] - g_i[j]$  from  $f(m[j] - g_i[j])$ :

$$\mathcal{L}(X_c; f, k) = \sum_j ((g_f[j] - g_i[j]) - f(m[j] - g_i[j]))^2 \quad (10)$$

For a given  $k$ , the best-fit polynomial minimizes this square-loss:

$$f_k^* = \arg \min_f \mathcal{L}(X_c; f, k) \quad (11)$$

To search over the space of polynomial models, we apply a well-studied technique called the Bayesian Information Criterion (BIC) [?]. This penalty can be interpreted as bias towards lower degree models, in other words, an Occam's Razor prior belief. This we reformulate the optimization problem in the following way to incorporate the BIC:

$$\arg \min_{f,k} |X_c| \log(\mathcal{L}(X_c; f, k)) + k \log(|X_c|) \quad (12)$$

The resulting optimal polynomial will tell how the regression affects varies as a function of  $m[j] - g_i[j]$  while controlling for over-fitting to our data. This optimization problem is non-convex so we incrementally try polynomials of degree 1,2,3.. etc. until we reach a local minimum.

## 6.2 Concentration of Grades

We can further estimate how much more concentrated changed grades are around the observed median. Recall in Section 4, we tested the significance of the absolute deviations using a Wilcoxon test statistic. The Wilcoxon statistic can be inverted to estimate a most likely *shift parameter*, that a constant shift  $\Delta$  in the distribution of absolute deviations  $X_c$  that maximally aligns them with  $X_n$  (ie.  $X_c + \Delta$  is most supported by the null hypothesis). Since  $X_c$  is a set of absolute deviations,  $\Delta$  tells us how much more concentrated  $X_c$  is than  $X_n$  around the observed medians. This parameter is relevant to the design of recommendation algorithms use proximity (eg. clustering or nearest neighbors).

We refer to [?] on the derivation of  $\Delta$  and its confidence interval:

$$D = \{x_n[j] - x_c[i]\} \forall i, j \in X_n, X_c \quad (13)$$

$$\Delta = \text{median}(D) \quad (14)$$

## 7. RESULTS

### 7.1 Observed Regression Towards the Median

In Figure 4, we plot  $m_i - g_i$  (the observed difference) against  $g_f - g_i$  (change in grade) for those participants that changed their grades. Supporting our initial hypothesis, we find that the values are positively correlated, which we define as a change towards the median.

Issue	N	Corr	P-Value
Obamacare	223	0.4580	5.1270e-20
K12	172	0.4813	1.6573e-18
College	139	0.4263	3.9772e-13
Immigration	105	0.5856	1.8984e-20
Marijuana	118	0.5397	3.0882e-19
Marriage Rights	105	0.5538	4.0921e-25

Furthermore, the correlations are significant with respect to the null correlation hypothesis. The significance test shows

it is highly unlikely that there is no correlation between the observed difference and the change in grade. Note that we discussed that this correlation does not on its own imply a tendency to regress towards the median grade as discussed in Section 4, as other models could result in similar correlations.

### 7.2 Non-Parametric Test Of Distance From the Median

We applied the non-parametric test proposed in Section 4. Figure 9 shows the mean absolute deviation for each group, and Table 7.2 tests its significance.




Figure 4: TODO

Issue	P ( $X_c$ vs. $X_n$ )	P ( $X'_c$ vs. $X_c$ )
Obamacare	0.0286	0.0161
K12	2.1314e-06	0.0086
College	1.3033e-04	0.0415
Immigration	7.3456e-07	4.4170e-05
Marijuana	2.7549e-10	4.2560e-05
Marriage Rights	3.5946e-06	2.4644e-10

For all of the issues, we find that set of absolute deviations from the median  $X_c$  is statistically significantly smaller compared to both  $X_n$  and  $X'_c$ . This suggests that the participants that changed their grades tended to be more tightly centered around median grade.

### 7.3 Sequence Dependence

Using the model proposed in Section 5, we calculated the test statistics for both the CRC and the Reference Survey. We found that for all issues the statistic was higher for the CRC suggesting an effect corroborating results in other work such as [?]. However, none of the results passed a  $p < 0.05$  statistical significance test. We believe that these results suggest that there is some sequence dependence in the CRC, however, we cannot definitively conclude that from the current quantity of data.

### 7.4 Grade Change Model

In Figure 7 and Figure 8, we show the results of our model search and locally optimal model for each issue. We found for four out of the six issues, K12, College, Immigration,




Figure 5: TODO

and Marijuana, the model we found was linear. However, for Obamacare and Marriage Rights, we found that the relationship was quadratic.

Figure 8 illustrates the nature of the quadratic relationship, and we see heterogeneity between a positive regression towards the median and a negative regression. Participants who initially graded the state higher than the median had a more significant tendency to regress downwards. This result is interesting for a few reasons: (1) contrary to our initial expectations the relationship is largely linear and (2) non-linearities appear in the two issues that received the highest grades which also happen to be highly politicized issues. There are many possible explanations for this including non-response bias [?] or aversive response [?]; and we defer a more detailed analysis to future work.




Figure 8: TODO

## 7.5 Shift-parameter estimation

In Figure 9, we show the results for the  $\Delta$  parameter estimate from inverting our hypothesis test. We find that on average grades in our change group were about half a letter grade (ie. +, -) closer to the median than grades from participants who didn't change. This parameter is relevant to both recommender systems and prediction tools. In the sim-

plest case, if we were to build a grade prediction tool that simply predicted the median grade, we could get misleadingly low prediction error. Likewise, algorithms that rely on proximity such as clustering or k-nearest neighbors could be misled to create a single big cluster around the median, when in fact the clustering around the median may be due to a biasing tendency.

## 7.6 Comparison To Reference Survey

We applied our proposed non-parametric test to compare the absolute deviations in the group of participants who changed their grades in the CRC with results from a reference survey. For the reference survey, we calculated the absolute deviation around the median (which the participants were not shown). We found for all but one issue the grades from the CRC were statistically significantly closer to the median than ones from the reference survey.

Issue	Med(Ref)	Med(CRC)	p-val
Obamacare	B	B	0.0078
K12	C+	C	0.3563
College	C-	C-	0.0011
Immigration	C	C+	0.0277
Marijuana	C	C	0.0076
Marriage Rights	B+	B+	0.0494

Furthermore, the two surveys aligned nearly perfectly in aggregate.

## 8. FUTURE WORK

The methods we proposed have several interesting directions of future interest. We want to extend our work to quantify biases in textual data. The California Report Card collects textual suggestions from participants in addition to the quantitative assessment results. Participants are encouraged to read the responses of others before leaving a suggestion of their own. We suspect that this may lead to a bias in the topics discussed by participants, and we would like to explore how similar non-parametric models can be extended to textual data.

Another compelling direction is to attempt to parameterize our model. We will explore whether we can model the grades as a mixture of binomial distributions (a discrete analog of a mixture of gaussians), and try to derive optimal tests and models for this data. Intuitively, parametrization should lead to increased statistical power and better fitting models; assuming that the data fits the underlying parametrization.

## 9. CONCLUSION

We proposed non-parametric hypothesis tests and models to evaluate the biasing tendency of visible aggregate statistics in the California Report Card. We found that revealing the median led to a statistically significantly tighter grouping of grades around the shown median grade.

We modeled the biasing effect as a regression towards the median grade and fit polynomial to represent the functional relationship between a participant's observed difference with the median and then subsequent grade change. We applied an information theoretic criteria to select a model of appropriate complexity. We found that this relationship was

quadratic in two out of the six issues, representing a heterogeneity in biasing for positive and negative differences with the median. We further showed how non-parametric ideas could be extended to the problem of Wilcoxon shift parameter estimation and quantify the effects of the biasing tendency.

In principle, the methods we proposed can be applied to test and model biases in a wide variety input mechanisms. This is a key motivation for our non-parametric approach. Understanding these biases, can give insight into the behavior of recommender systems that train on such data.

## 10. ACKNOWLEDGMENTS

## 11. REFERENCES

- [1] J. Albors, J. C. Ramos, and J. L. Hervás. New learning network paradigms: Communities of objectives, crowdsourcing, wikis and open source. *International Journal of Information Management*, 28(3):194–202, 2008.
- [2] S. E. Asch. Opinions and social pressure. *Readings about the social animal*, pages 17–26, 1955.
- [3] S. E. Asch. *Studies of independence and conformity*. American Psychological Association, 1956.
- [4] A. V. Banerjee. A simple model of herd behavior. *The Quarterly Journal of Economics*, 107(3):797–817, 1992.
- [5] S. Bikhchandani and S. Sharma. Herd behavior in financial markets: A review. 2000.
- [6] R. Bond and P. B. Smith. Culture and conformity: A meta-analysis of studies using asch’s (1952b, 1956) line judgment task. *Psychological bulletin*, 119(1):111, 1996.
- [7] R. E. Burnkrant and A. Cousineau. Informational and normative social influence in buyer behavior. *Journal of Consumer research*, pages 206–215, 1975.
- [8] C. Danescu-Niculescu-Mizil, G. Kossinets, J. Kleinberg, and L. Lee. How opinions are received by online communities: a case study on amazon. com helpfulness votes. In *Proceedings of the 18th international conference on World wide web*, pages 141–150. ACM, 2009.
- [9] P. M. DeMarzo, D. Vayanos, and J. Zwiebel. Persuasion bias, social influence, and unidimensional opinions. *The Quarterly Journal of Economics*, 118(3):909–968, 2003.
- [10] U. M. Dholakia, S. Basuroy, and K. Soltysinski. Auction or agent (or both)? a study of moderators of the herding bias in digital auctions. *International Journal of Research in Marketing*, 19(2):115–130, 2002.
- [11] R. M. Groves, F. J. Fowler Jr, M. P. Couper, J. M. Lepkowski, E. Singer, and R. Tourangeau. *Survey methodology*. John Wiley & Sons, 2013.
- [12] J.-H. Huang and Y.-F. Chen. Herding in online product choice. *Psychology & Marketing*, 23(5):413–428, 2006.
- [13] E. L. Lehmann and H. J. D’Abrera. *Nonparametrics: statistical methods based on ranks*. Springer New York, 2006.
- [14] J. Lorenz, H. Rauhut, F. Schweitzer, and D. Helbing. How social influence can undermine the wisdom of crowd effect. *Proceedings of the National Academy of Sciences*, 108(22):9020–9025, 2011.
- [15] S. Moscovici and C. Faucheux. Social influence, conformity bias, and the study of active minorities. *Advances in experimental social psychology*, 6:149–202, 1972.
- [16] L. Muchnik, S. Aral, and S. J. Taylor. Social influence bias: A randomized experiment. *Science*, 341(6146):647–651, 2013.
- [17] B. S. Noveck. Wiki-government. *Democracy: A Journal of Ideas* (7), 2008.
- [18] K. O’Hara. Transparency, open data and trust in government: Shaping the infosphere. In *Proceedings of the 3rd Annual ACM Web Science Conference*, pages 223–232. ACM, 2012.
- [19] B. Shneiderman. *Designing the user interface: strategies for effective human-computer interaction*, volume 2. Addison-Wesley Reading, MA, 1992.
- [20] R. Sipos, A. Ghosh, and T. Joachims. Was this review helpful to you? it depends! context and voting patterns in online content.
- [21] W. Wood. Attitude change: Persuasion and social influence. *Annual review of psychology*, 51(1):539–570, 2000.
- [22] H. Zhu, B. Huberman, and Y. Luon. To switch or not to switch: understanding social influence in online choices. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, pages 2257–2266. ACM, 2012.

## APPENDIX

### A. HEADINGS IN APPENDICES




Figure 3: TODO

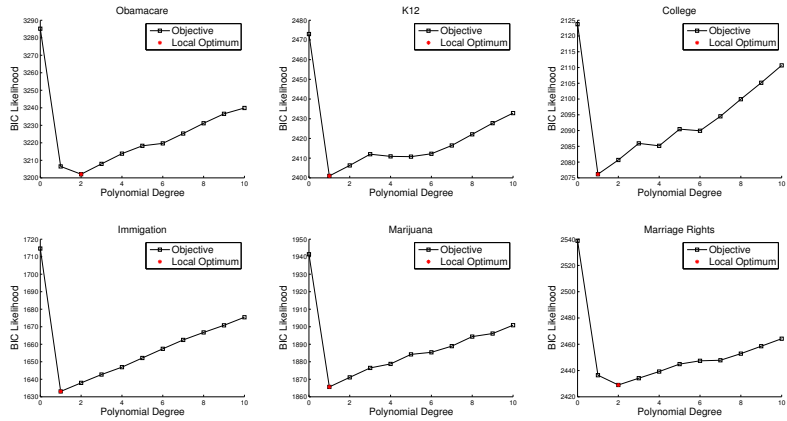


Figure 6: TODO




Figure 7: TODO